

Lecture 5: Evolutionary Computation

Artificial Intelligence
CS-UY-4613-A / CS-GY-6613-I

Julian Togelius
julian.togelius@nyu.edu

Table of contents

- Recap: hill-climbing, search as optimization
- Biological metaphor
- Recombination, mutation, selection
- Representations
- Fitness functions
- Constraints
- Multiple objectives
- Coevolution

Local search algorithms

- In many optimization problems, the path to the goal is irrelevant; the goal state itself is the solution
- State space = set of "complete" configurations
- Find configuration satisfying constraints, e.g., n-queens
- In such cases, we can use local search algorithms: keep a single "current" state, try to improve it

Hill-climbing

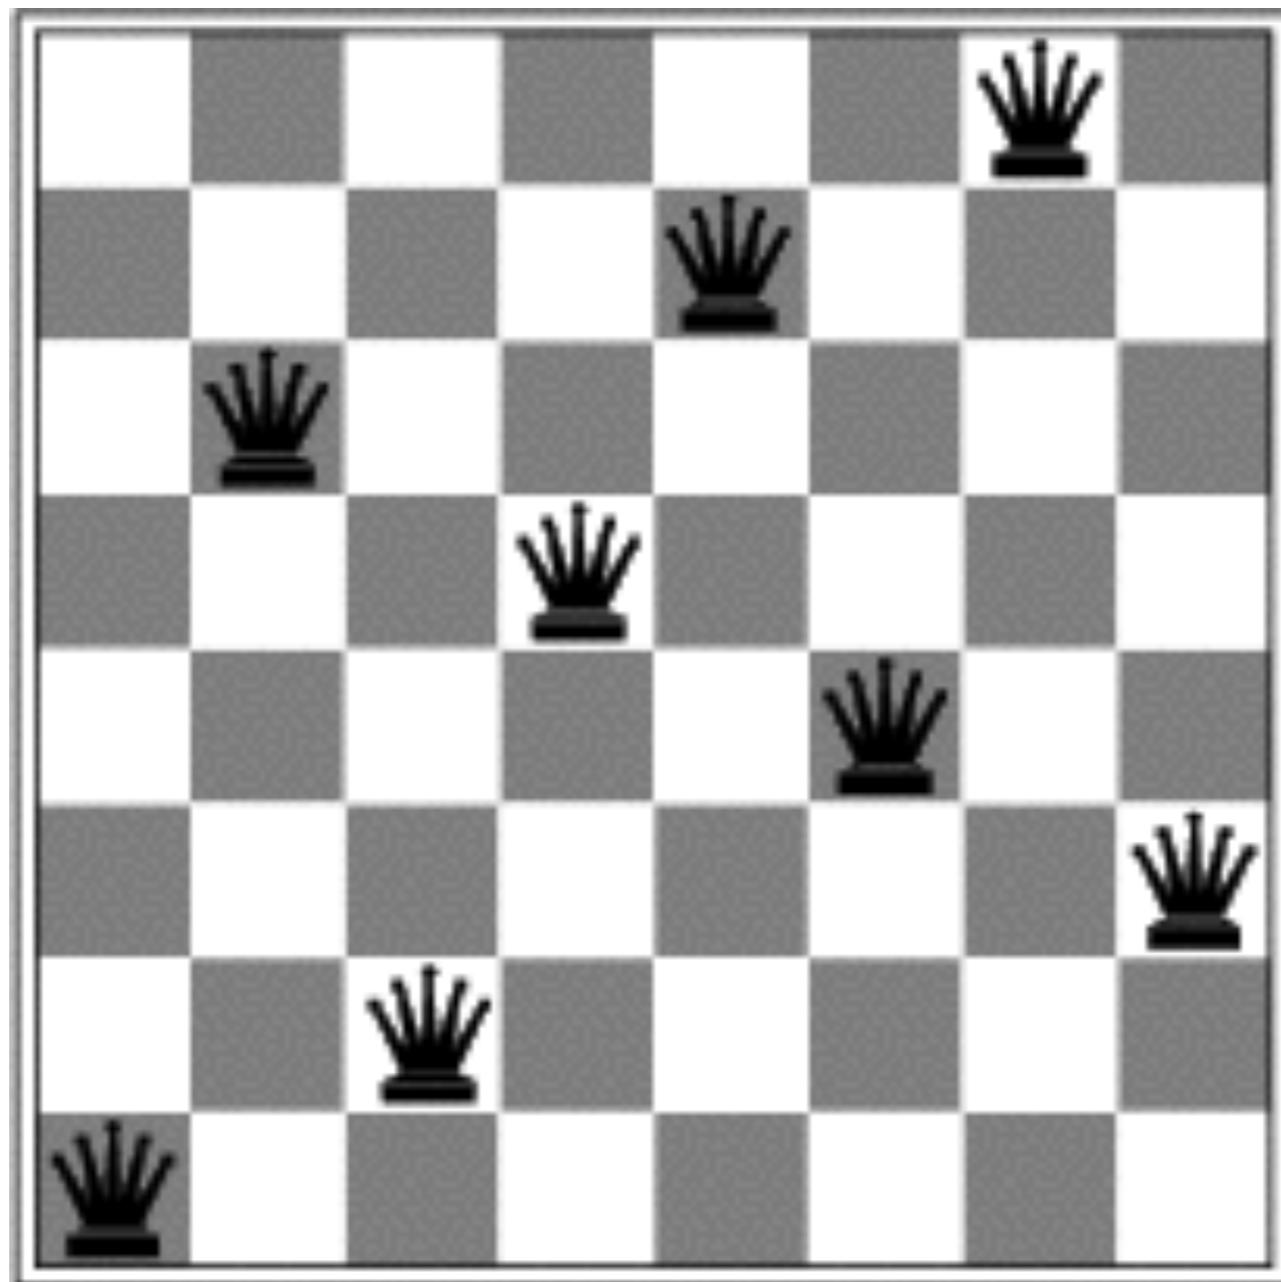
```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                  neighbor, a node
  current  $\leftarrow$  MAKE-NODE(INITIAL-STATE[problem])
  loop do
    neighbor  $\leftarrow$  a highest-valued successor of current
    if VALUE[neighbor]  $\leq$  VALUE[current] then return STATE[current]
    current  $\leftarrow$  neighbor
```

n-queens

- Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal



Local minimum



Other ideas...

- Exhaustive search: try all configurations, one after another
- Random search: generate random configurations

Simulated annealing

- Do bad moves with decreasing probability

```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state
    inputs: problem, a problem
            schedule, a mapping from time to "temperature"
    local variables: current, a node
                    next, a node
                    T, a "temperature" controlling prob. of downward steps
    current  $\leftarrow$  MAKE-NODE(INITIAL-STATE[problem])
    for t  $\leftarrow$  1 to  $\infty$  do
        T  $\leftarrow$  schedule[t]
        if T = 0 then return current
        next  $\leftarrow$  a randomly selected successor of current
         $\Delta E \leftarrow$  VALUE[next] - VALUE[current]
        if  $\Delta E > 0$  then current  $\leftarrow$  next
        else current  $\leftarrow$  next only with probability  $e^{\Delta E/T}$ 
```

Local beam search

- Keep track of k states rather than just one
- Start with k randomly generated states
- At each iteration, all the successors of all k states are generated
- If any one is a goal state, stop; else select the k best successors from the complete list and repeat.

Limitations of local search

- Susceptible to local optima
- No sharing of information between parallel hill-climbing runs
- No possibility to learn from previous search experience

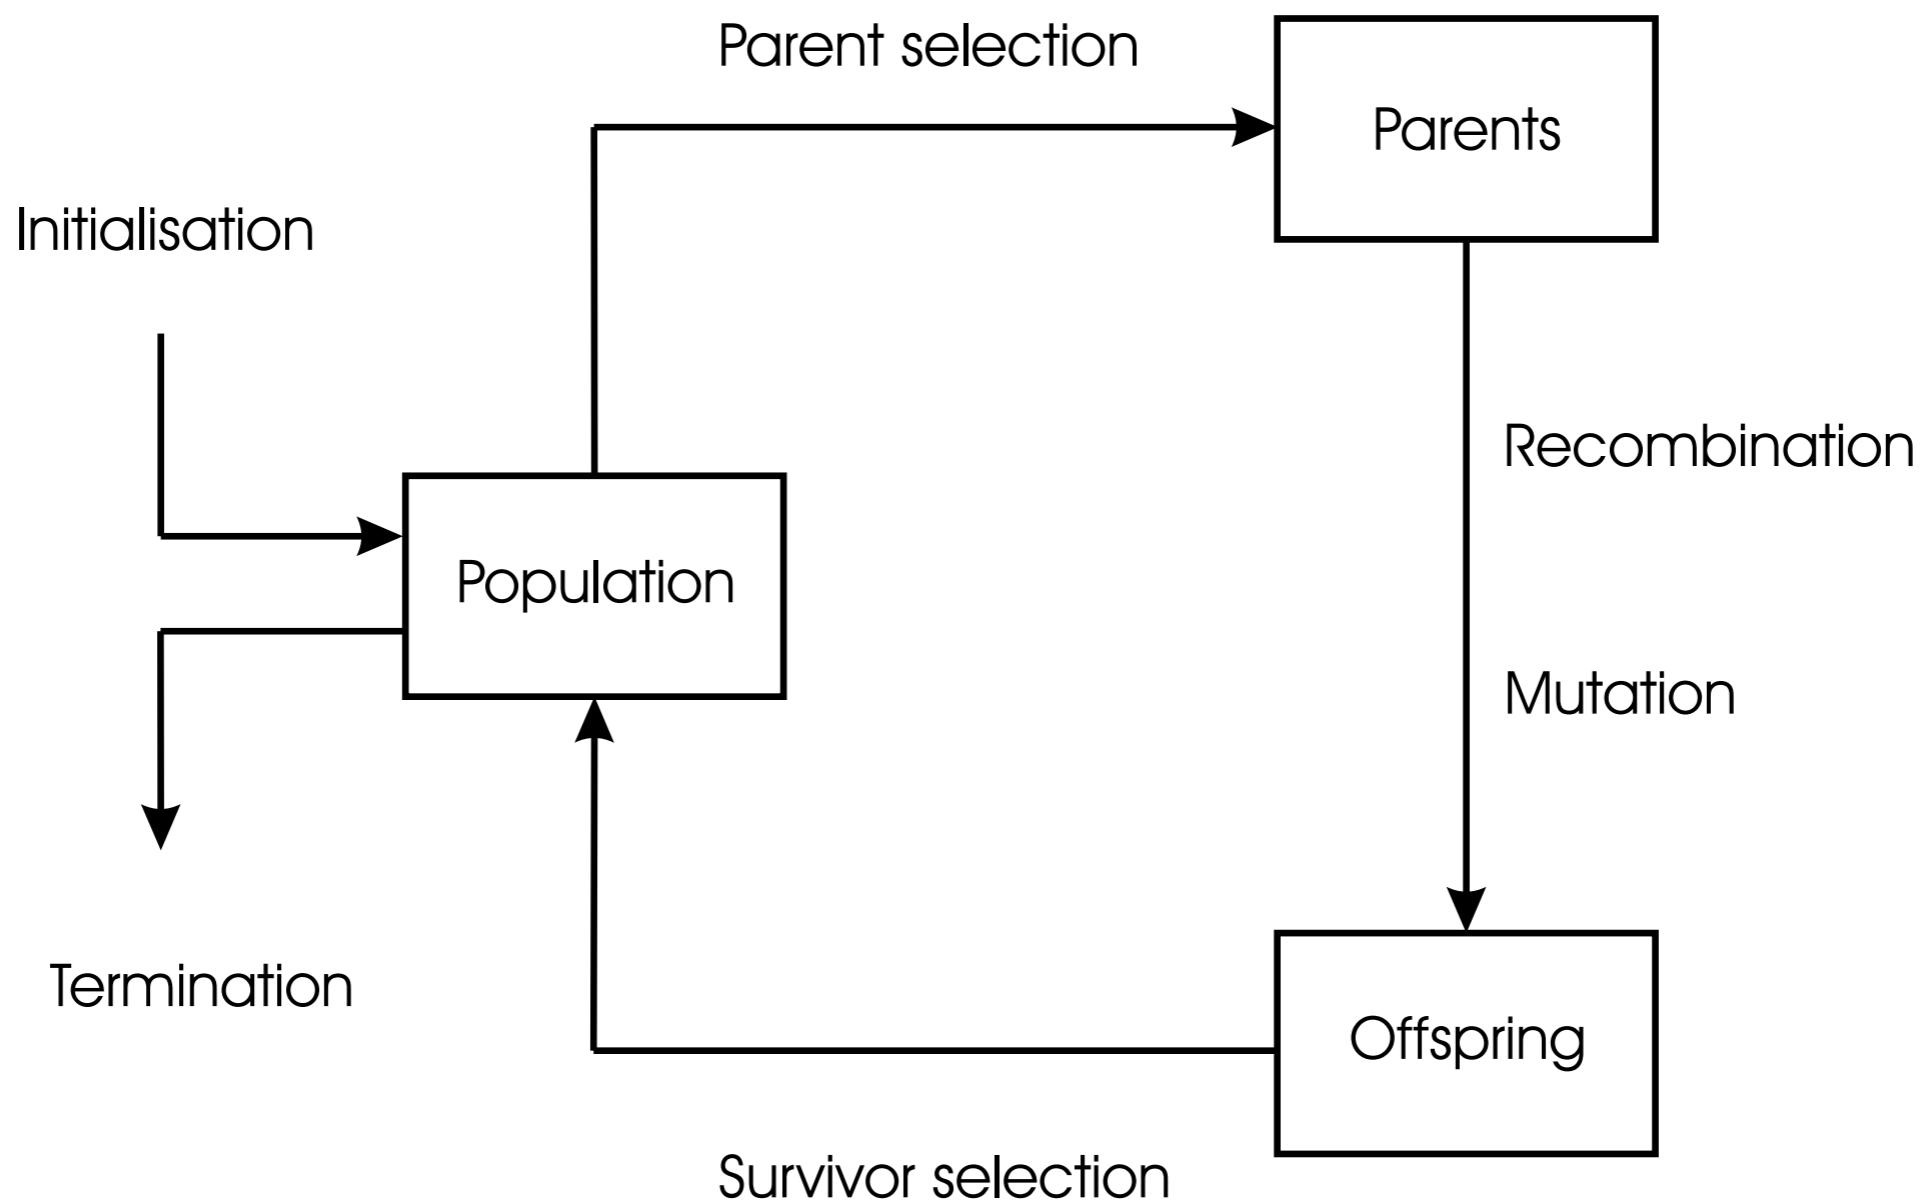
Evolutionary algorithms

- Stochastic global optimisation algorithms
- Inspired by Darwinian natural evolution
- Extremely domain-general, widely used in practice



© Reuters

Generic EA



What's essential?

- Fitness evaluation
- Selection
 - Fitness influences how many offspring an individual has
- Variation
 - Could be mutation *and/or* crossover

Generic EA

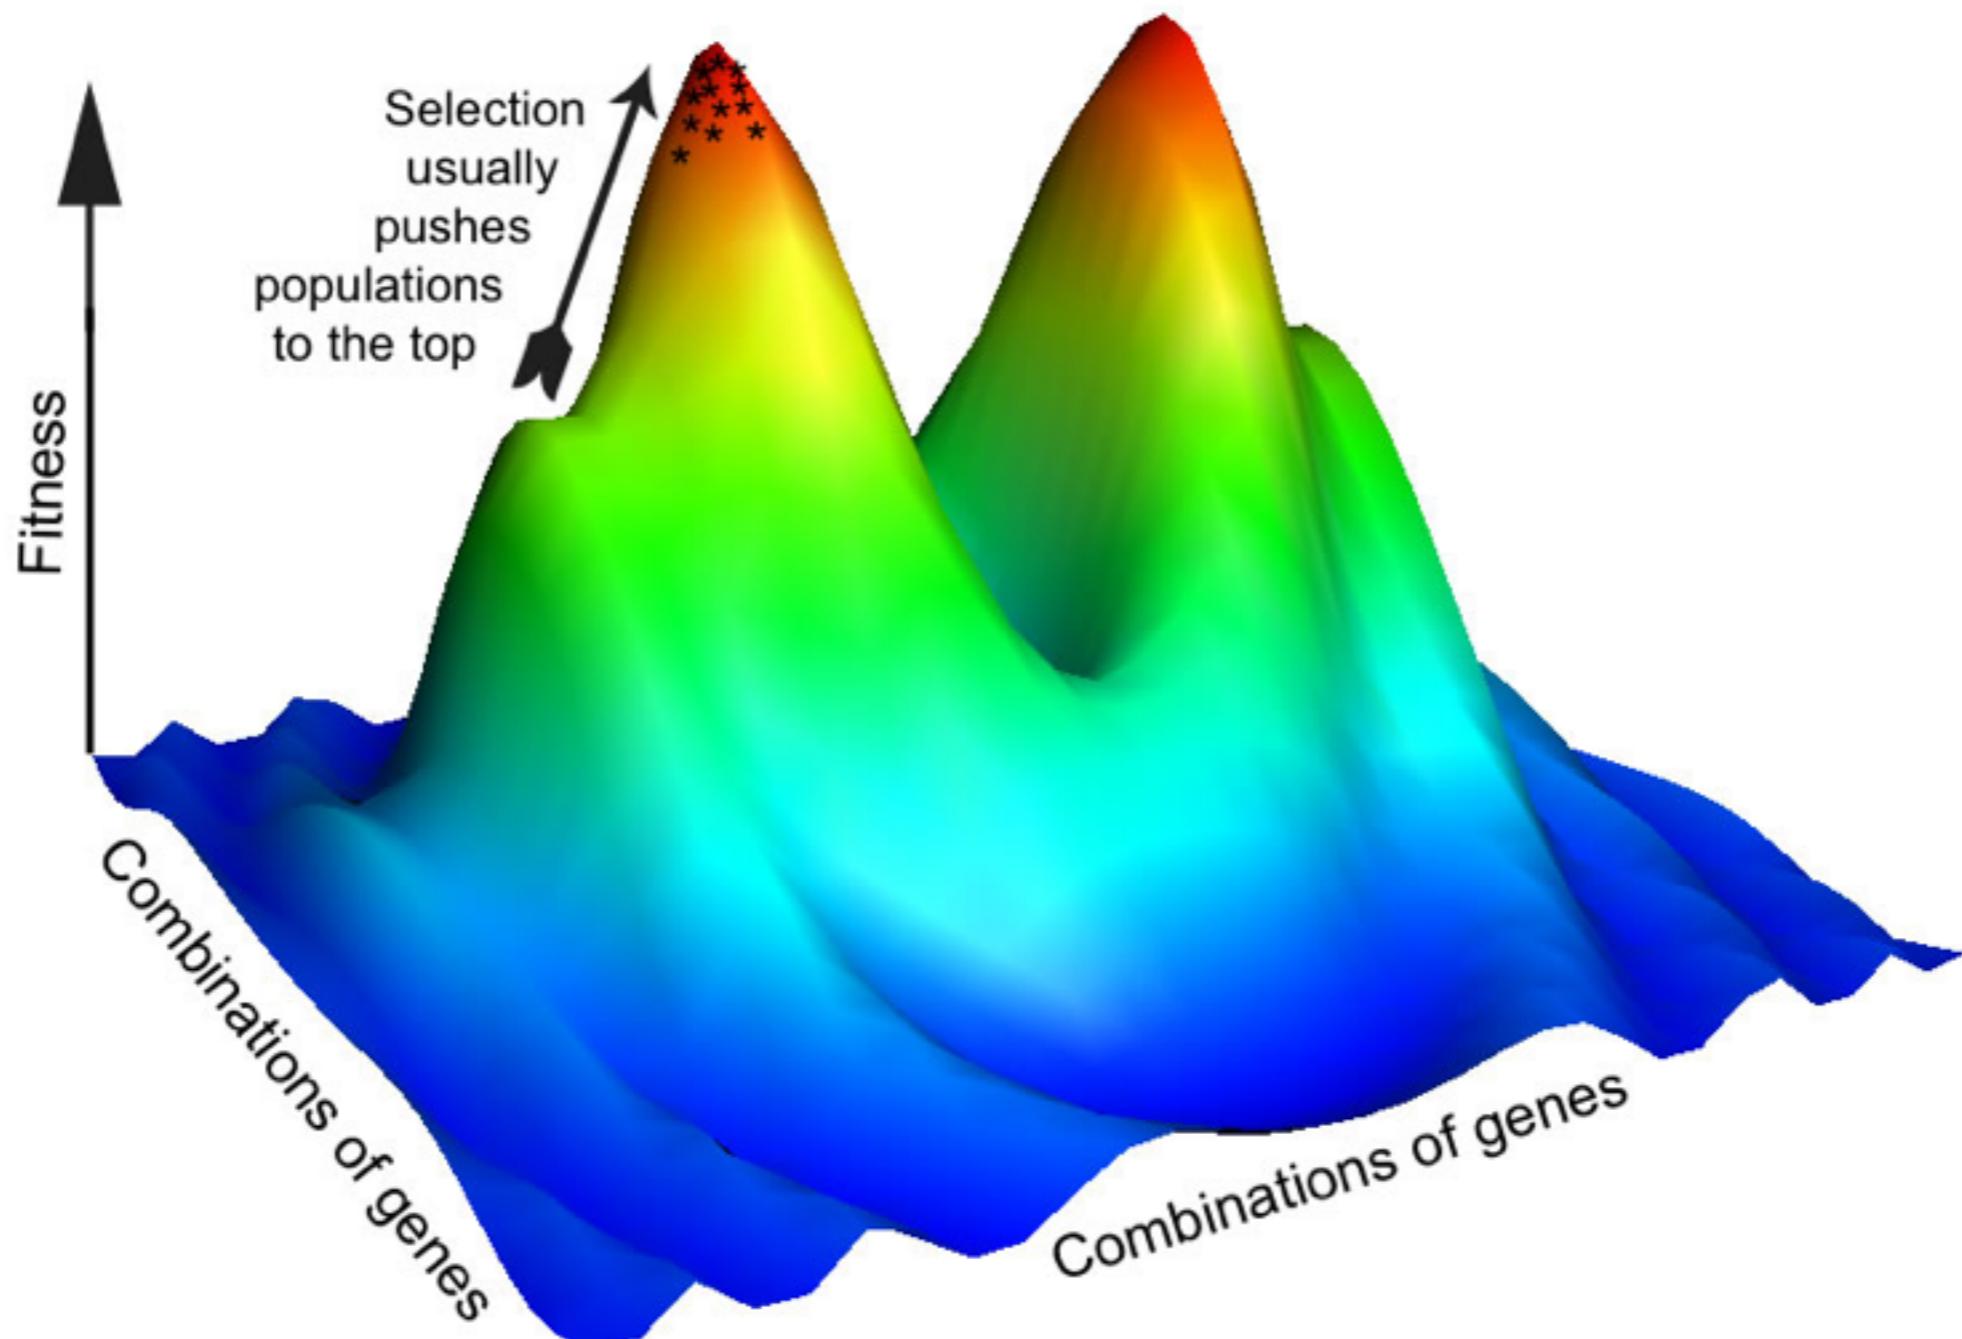
BEGIN

INITIALISE population with random candidate solutions;
EVALUATE each candidate;
REPEAT UNTIL (*TERMINATION CONDITION* is satisfied) DO
 1 *SELECT* parents;
 2 *RECOMBINE* pairs of parents;
 3 *MUTATE* the resulting offspring;
 4 *EVALUATE* new candidates;
 5 *SELECT* individuals for the next generation;

OD

END

The fitness landscape



Simple $\mu+\lambda$ Evolution Strategy

- Create a population of $\mu+\lambda$ individuals
- Each generation
 - Evaluate all individuals in the population
 - Sort by fitness
 - Remove the worst λ individuals
 - Replace with mutated copies of the μ best

Example: OneMax

- Representation: binary strings of length 5
- Examples: 00000, 01001, 00100, 11101, 11111
- Fitness function: count the number of ones
- Mutation: flip a random bit

Example: OneMax

- Initial population: 00100, 11001, 10010, 01010
- Gen 1: 11001 : 3, 10010 : 2, 01010 : 2, 00100 : 1
Selection: 11001 : 3, 10010 : 2
- Gen 2: 11101 : 4, 11001 : 3, 10010 : 2, 00010 : 1
Selection: 11101 : 4, 11001 : 3
- Gen 3: 11111 : 5, 11101 : 4, 11101 : 4, 11001 : 3

Generally, you need

- A solution representation
- Variation operators (mutation and/or crossover)
- A fitness (evaluation) function

Representation

- Individuals in the population are referred to as genotypes (e.g. arrays of doubles)
 - All variation and replication mechanisms are performed on the genotype
- What is evaluated by the fitness function is the phenotype (e.g. neural networks)
- Genotype-to-phenotype mapping should be fast, deterministic, preserve locality

Representations should be...

- Compact
- Expressive (represent all “interesting” parts of phenotype space)
- Human-understandable (?)
- Local (?)

Evaluation function

- Also called fitness function
- Returns a number (or ranking) of the evaluated individual (phenotype)
- You could have one or many
- Evaluation functions should be smooth, fast to evaluate, noise-free, accurate and relevant

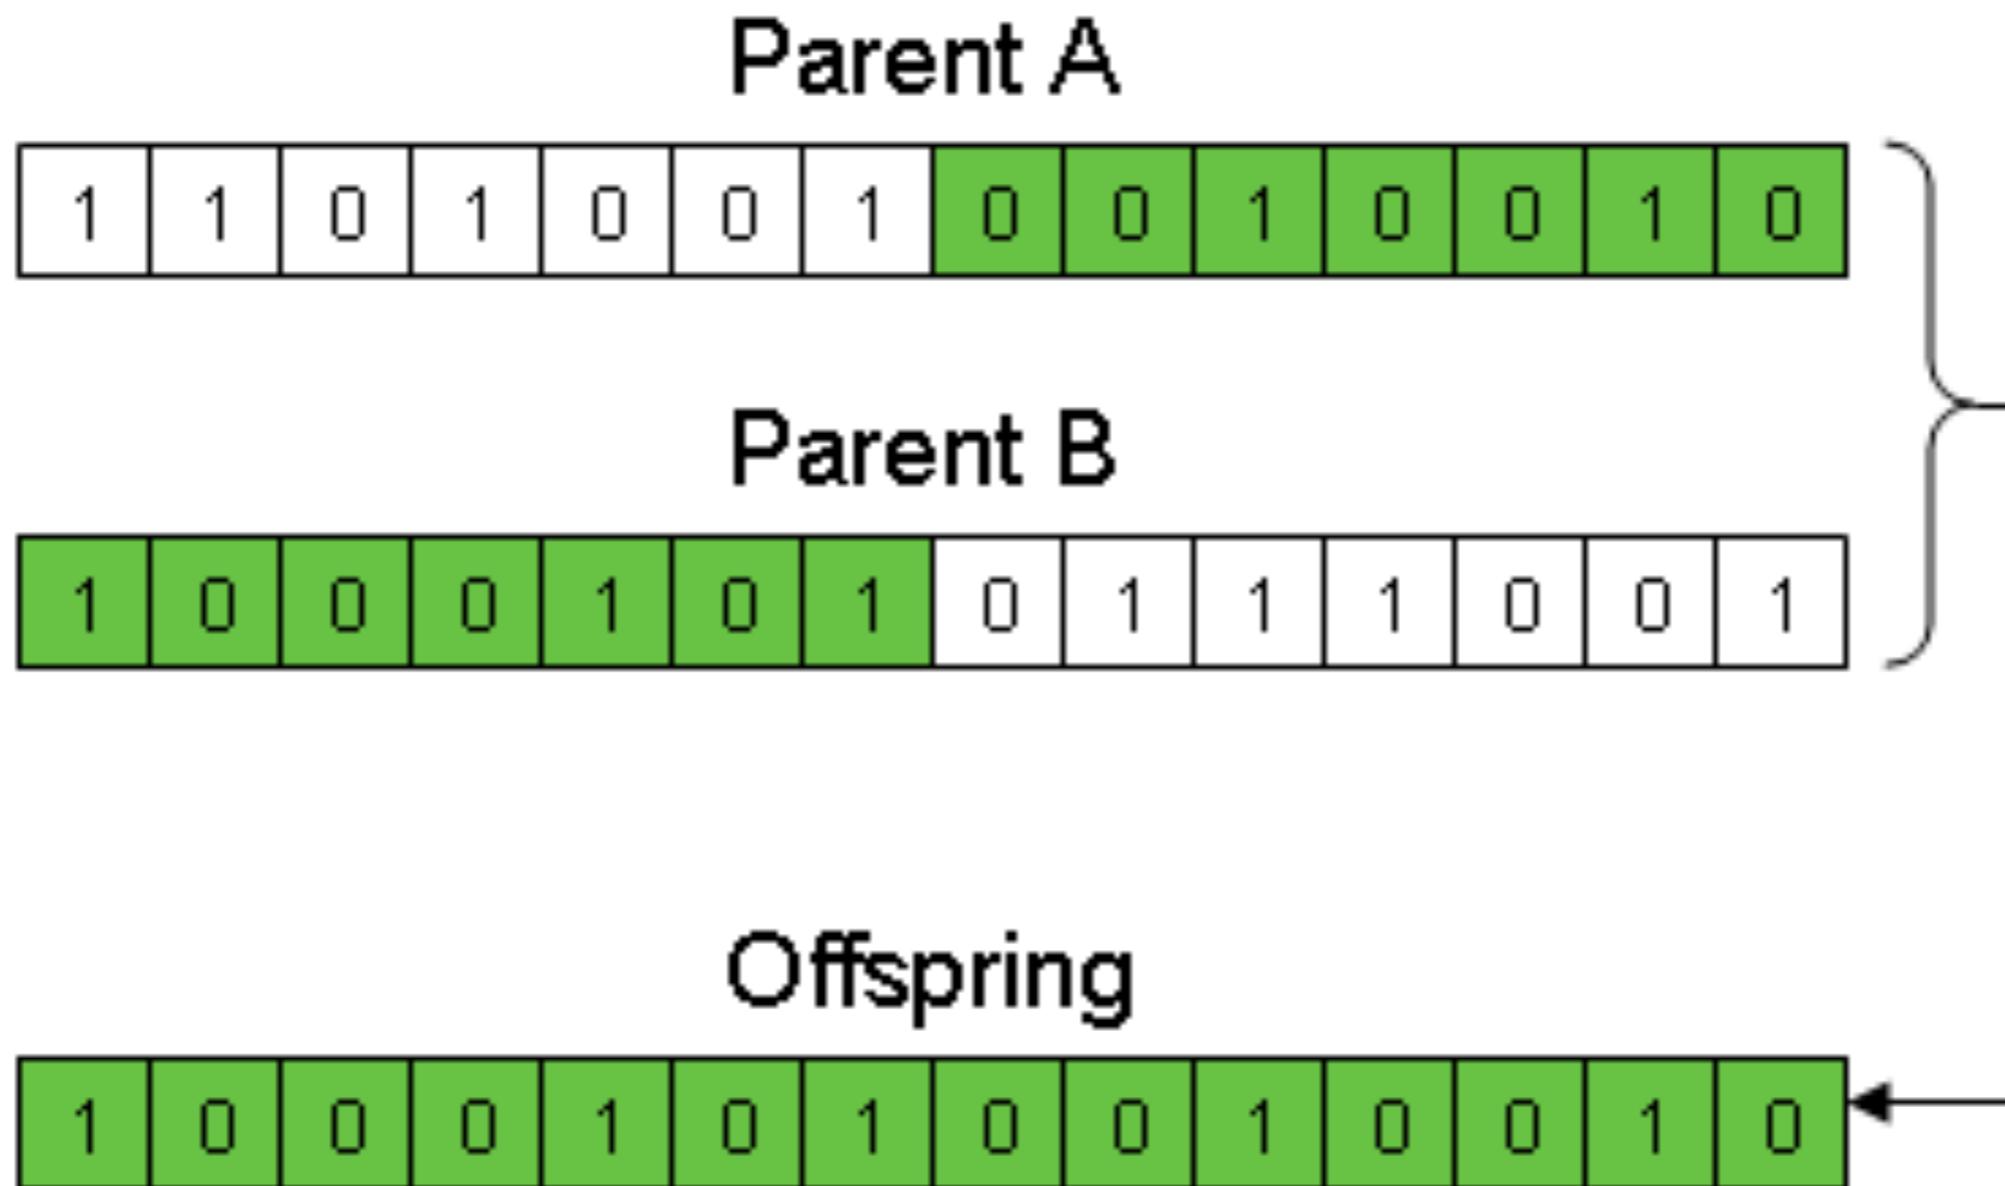
Parent selection

- Which individuals get to reproduce?
- Can choose this based on fitness...
- ...or fitness ranking
- Deterministic or stochastic
- How many to select?

Variation operators

- Should be appropriate to the representation
- Mutation
 - Should have the appropriate magnitude
 - e.g. Gaussian mutation: add values drawn from Gaussian distribution to an array of real numbers. Gaussian distribution is a random distribution with a mean and a standard deviation. e.g. in Java: `random.nextGaussian()`
 - Or: add a value to all parts of the genome drawn from uniform distribution. e.g. `Math.random() * some small value`
- Recombination (crossover)
 - Only do this if you know what you are doing

Example: one-point crossover



Survivor selection

- Elitism?
 - Usually a good thing
- Who to replace?
 - Always kill the worst?
 - Diversity preservation?

Initialization

- Population should be diverse enough
- Should you start from results of previous evolutionary runs?

Termination condition

- When you're done?
- When you've waited long enough?
- Relevant factor to count: number of evaluations

Simple $\mu+\lambda$ ES (concrete)

- Create a population of $\mu+\lambda$ individuals (e.g. 50+50) represented as e.g. vectors of real numbers in [0..1]
- Each generation (until 100 generations)
 - Evaluate all individuals in the population
 - Sort by fitness (e.g. win rate or score for the game; higher is better)
 - Remove the worst λ individuals
 - Replace with mutated copies of the μ best (mutate through Gaussian mutation with mean 0 and s.d. 0.1)

Types of evolutionary algorithms

- Genetic Algorithms
 - Representation: bitstring
 - Mostly driven by crossover
- Evolution Strategies
 - Representation: real values
 - Mostly driven by (self-adaptive) mutation
- Genetic Programming
 - Representation: expression trees
- Estimation of Distribution Algorithms, e.g. CMA-ES
- Others: Particle Swarm Optimization, Differential Evolution...

```

function GENETIC-ALGORITHM(population, FITNESS-FN) returns an individual
  inputs: population, a set of individuals
          FITNESS-FN, a function that measures the fitness of an individual

  repeat
    new-population  $\leftarrow$  empty set
    for i = 1 to SIZE(population) do
      x  $\leftarrow$  RANDOM-SELECTION(population, FITNESS-FN)
      y  $\leftarrow$  RANDOM-SELECTION(population, FITNESS-FN)
      child  $\leftarrow$  REPRODUCE(x, y)
      if (small random probability) then child  $\leftarrow$  MUTATE(child)
      add child to new-population
    population  $\leftarrow$  new-population
  until some individual is fit enough, or enough time has elapsed
  return the best individual in population, according to FITNESS-FN

```

```

function REPRODUCE(x, y) returns an individual
  inputs: x, y, parent individuals

  n  $\leftarrow$  LENGTH(x); c  $\leftarrow$  random number from 1 to n
  return APPEND(SUBSTRING(x, 1, c), SUBSTRING(y, c + 1, n))

```

Figure 4.8 A genetic algorithm. The algorithm is the same as the one diagrammed in Figure ??, with one variation: in this more popular version, each mating of two parents produces only one offspring, not two.

Taking a step back...
How would you actually
use an evolutionary
algorithm?



Uses of evolution:

- Evolve a route
- Evolve the parameters of the car
- Evolve a controller that can drive the car

MARIO: 1024

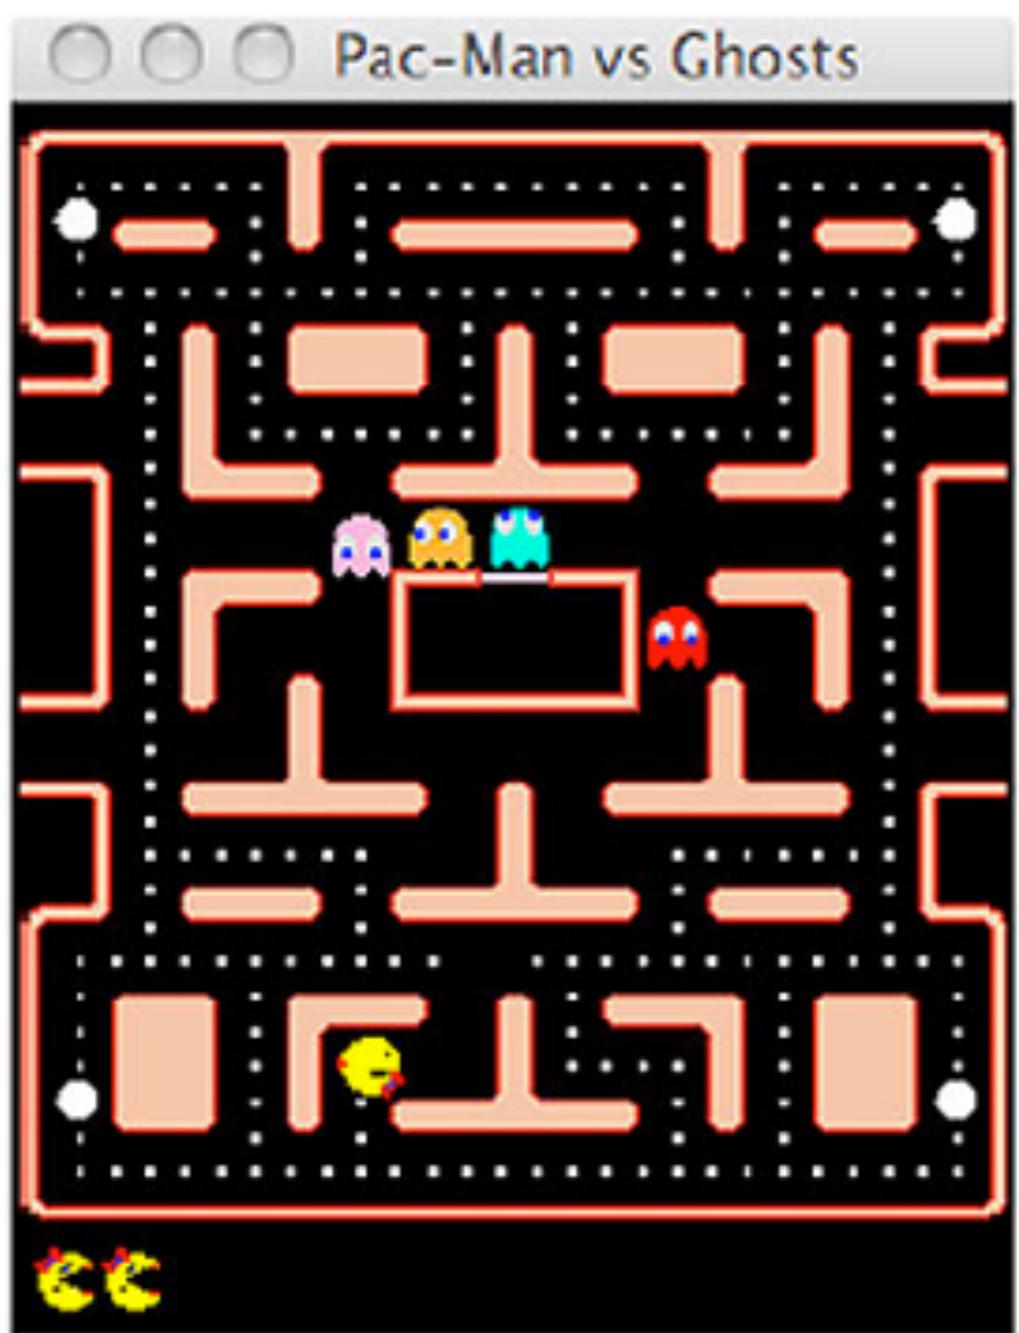
COINS
18

DIFFICULTY 10
TIME 144
WorldPause
False



Uses of evolution:

- Evolve level, heuristic (for search algorithm), search algorithm, path, physics, physics prediction, enemies....



Uses of evolution:

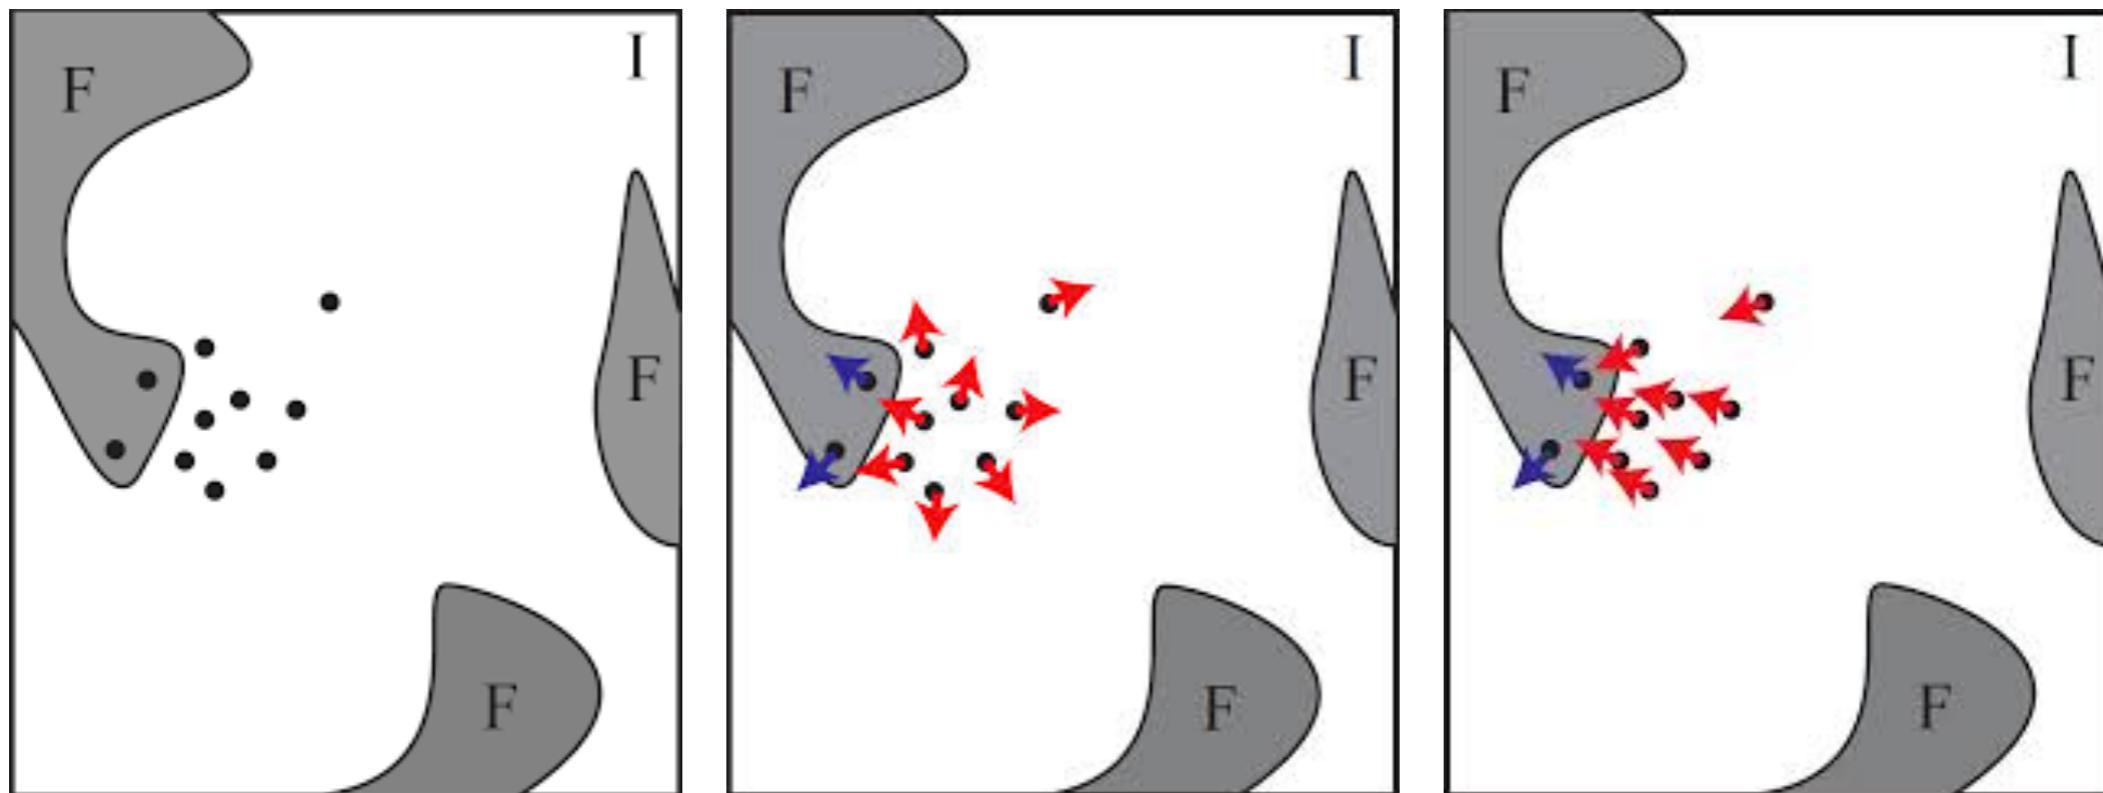


Uses of evolution:

Evolution with constraints

- Sometimes there are certain features which *must* be part of the solution
 - e.g. car has four wheels, path passes through New Jersey, wing length less than 40m
- These could be encoded in the fitness function - but gets complicated
- Better, encode them as constraints

Feasible and Infeasible Areas



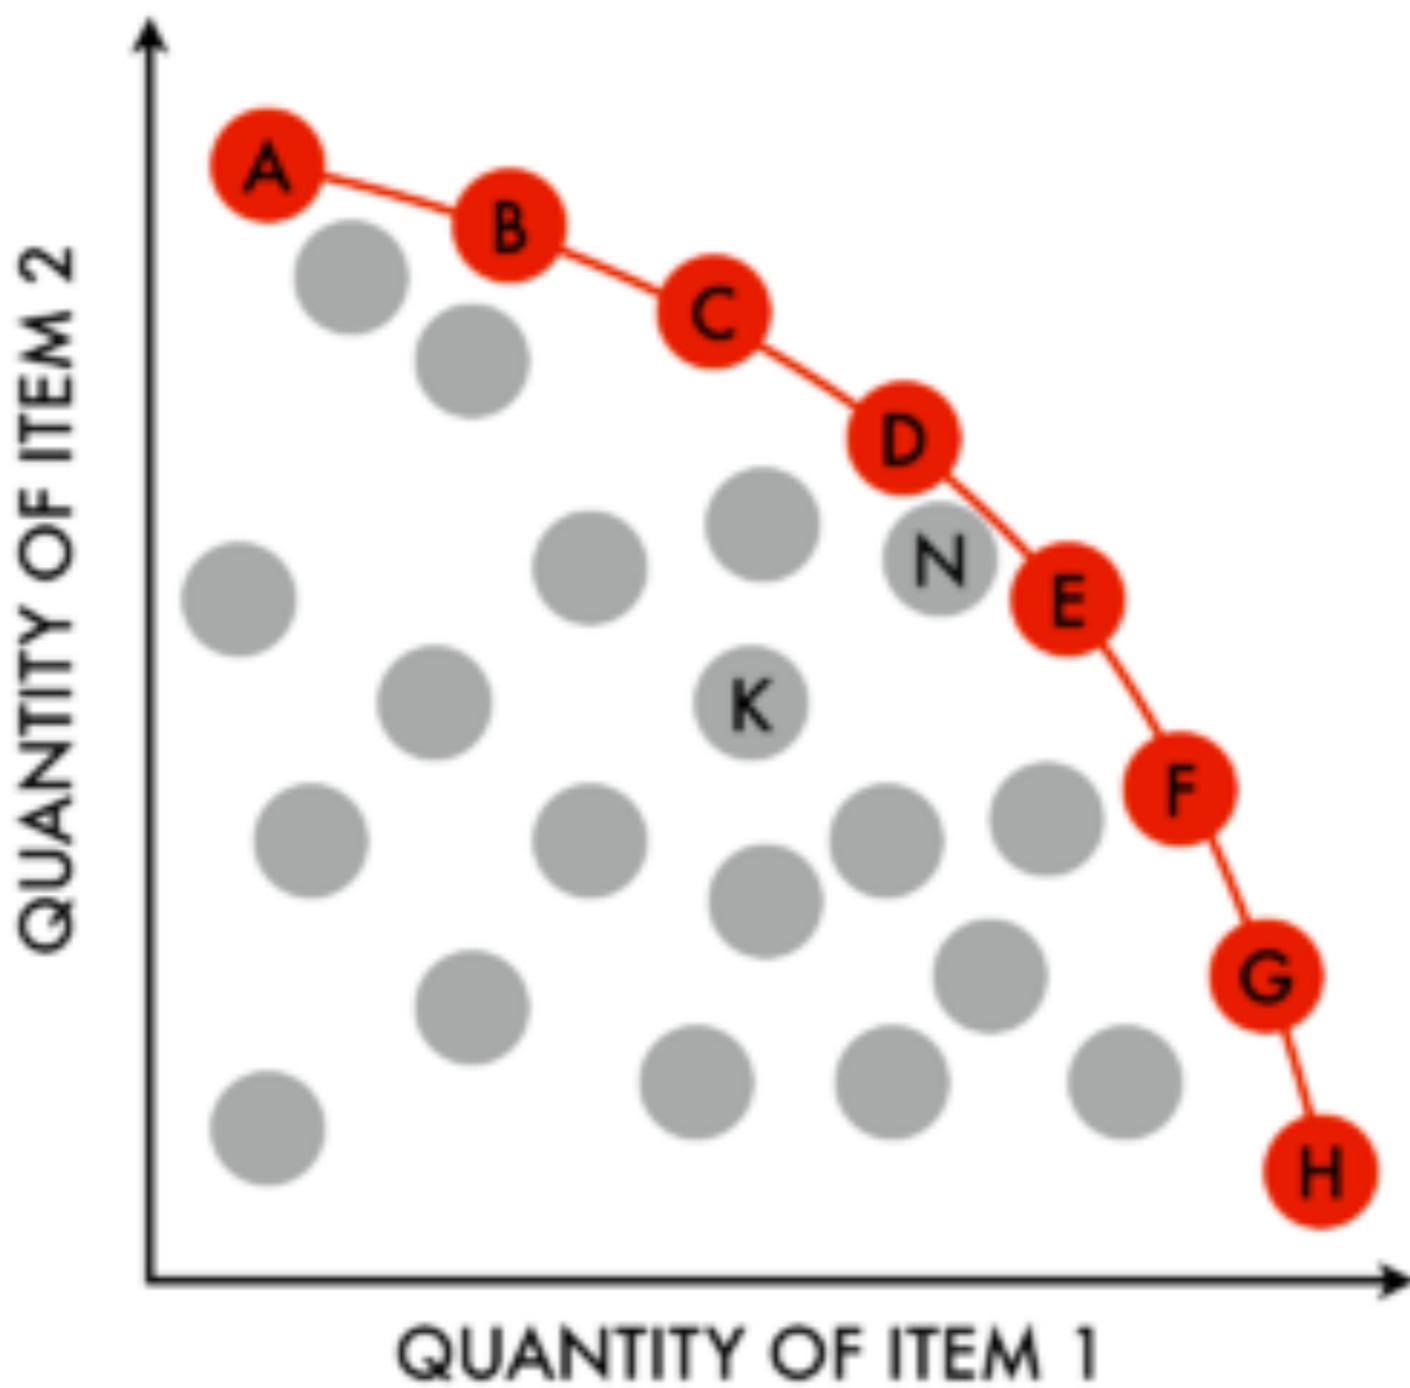
FI-2Pop algorithm

- Keep two populations: feasible and infeasible individuals
- Feasible individuals are selected for their fitness
- Infeasible individuals are selected for their proximity to feasibility
- Infeasible offspring of feasible parents are moved to infeasible population

Multiobjective evolution

- Often, you want to optimize for multiple fitness dimensions
 - For example, you want the car to be fuel-efficient, cheap and fast...
- These objectives are often *partially conflicting*
- A multiobjective evolutionary algorithm finds the *pareto front* of solutions given these objectives

Pareto fronts





© Reuters



Competitive coevolution

- Two (or more populations) with different fitness functions
 - e.g. predators and prey
- The fitness of one population depends on the other and vice versa
- Goal: start an *arms race*

