

Lecture 6: Adversarial Search

Artificial Intelligence
CS-UY-4613-A / CS-GY-6613-I
Julian Togelius
julian.togelius@nyu.edu

Extra! Hear all about it!

- Adversarial problems and games
- Minimax
- Alpha-beta pruning
- Heuristics and considerations
- After the break: Algorithms to implement

Adversarial problems?

- So far, problems have been deterministic with a known forward model
 - Though the methods would work with noise and/or imperfect models
- Adversarial: there is another agent, which works against you
 - The opponent is not predictable
- Game trees rather than search trees


Paradigm case



Other cases




Two-player game tree



Minimax

- Perfect play for deterministic games
- Choose move to state with highest minimax value (best payoff against perfect opponent)



function MINIMAX-DECISION($state$) **returns** *an action*

$v \leftarrow \text{MAX-VALUE}(state)$

return the *action* in SUCCESSORS($state$) with value v

function MAX-VALUE($state$) **returns** *a utility value*

if TERMINAL-TEST($state$) **then return** UTILITY($state$)

$v \leftarrow -\infty$

for a, s in SUCCESSORS($state$) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$

return v

function MIN-VALUE($state$) **returns** *a utility value*

if TERMINAL-TEST($state$) **then return** UTILITY($state$)

$v \leftarrow \infty$

for a, s in SUCCESSORS($state$) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$

return v


```

MinMax (GamePosition game) {
    return MaxMove (game);
}

MaxMove (GamePosition game) {
    if (GameEnded(game)) {
        return EvalGameState(game);
    }
    else {
        best_move <- {};
        moves <- GenerateMoves(game);
        ForEach moves {
            move <- MinMove(ApplyMove(game));
            if (Value(move) > Value(best_move)) {
                best_move <- move;
            }
        }
        return best_move;
    }
}

MinMove (GamePosition game) {
    best_move <- {};
    moves <- GenerateMoves(game);
    ForEach moves {
        move <- MaxMove(ApplyMove(game));
        if (Value(move) > Value(best_move)) {
            best_move <- move;
        }
    }

    return best_move;
}

```

Properties of Minimax


- *Complete?* Yes (if tree is finite)
- *Optimal?* Yes (against an optimal opponent)
- *Time complexity?* $O(b^m)$
- *Space complexity?* $O(bm)$ (depth-first exploration)

For chess, $b \approx 35$, $m \approx 100$ for "reasonable" games
> exact solution completely infeasible


α - β pruning

- Can we improve this?
- Idea: don't consider branches of the tree that cannot lead to a better outcome than those that we have already explored


α - β pruning




α - β pruning




α - β pruning



α - β pruning



α - β pruning




Properties of α - β

- Pruning does **not** affect final result
- Good move ordering improves effectiveness of pruning
- With "perfect ordering", time complexity = $O(b^{m/2})$
 - doubles depth of search
- A simple example of the value of reasoning about which computations are relevant (metareasoning)

Why is it called α - β ?

- α is the value of the best (i.e. highest-value) choice found so far at any choice point along the path for *max*
- If v is worse than α , *max* will avoid it
> prune that branch
- Define β similarly for *min*



function ALPHA-BETA-SEARCH(*state*) **returns** *an action*

inputs: *state*, current state in game

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the *action* in SUCCESSORS(*state*) with value v

function MAX-VALUE(*state*, α , β) **returns** *a utility value*

inputs: *state*, current state in game

α , the value of the best alternative for MAX along the path to *state*

β , the value of the best alternative for MIN along the path to *state*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for a, s in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s, \alpha, \beta))$

if $v \geq \beta$ **then return** v

$\alpha \leftarrow \text{MAX}(\alpha, v)$

return v

```

function MIN-VALUE( $state, \alpha, \beta$ ) returns a utility value
  inputs:  $state$ , current state in game
            $\alpha$ , the value of the best alternative for MAX along the path to  $state$ 
            $\beta$ , the value of the best alternative for MIN along the path to  $state$ 

  if TERMINAL-TEST( $state$ ) then return UTILITY( $state$ )
   $v \leftarrow +\infty$ 
  for  $a, s$  in SUCCESSORS( $state$ ) do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s, \alpha, \beta))$ 
    if  $v \leq \alpha$  then return  $v$ 
     $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return  $v$ 

```


Still, there is no time...

- Suppose we have 100 secs, explore 10^4 nodes/sec
> 10^6 nodes per move
 - a far cry from 35^{100} ...
- Standard approach: use an *evaluation function* (heuristic)
- Either at the same depth for all branches, or use a *cutoff test* such as in quiescence search

Evaluation functions

- Simplest: number of white pieces - number of black pieces
- More complex: assign values to piece types
- Even more complex: count number of threats, try to recognize known positions
- Generally: Linear weighted sum of features
- Neural network
- Read more: Blondie24 by David Fogel

Piece weights?



Cutting off search

- MinimaxCutoff is identical to MinimaxValue except
 - Terminal? is replaced by Cutoff?
 - Utility is replaced by Eval
- Does it work in practice?
 - $b^m = 10^6$, $b=35$ means $m=4$
- 4-ply lookahead is (in general) a hopeless chess player!
 - 4-ply \approx human novice
 - 8-ply \approx old-school Chess program, human master
 - 12-ply \approx Deep Blue, Kasparov

Some deterministic two-player games

- Chess: Kasparov vs Deep Blue 1997
- Checkers: Chinook almost defeated Tinsley (grand champion) 1994
 - Solved 2007
- Othello: computers vastly better than humans
- Go: computers currently at advanced intermediate level, thanks to MCTS

Discussion on projects

Important

- ***GAME TREE SEARCH: YOU ARE SEARCHING IN STATE SPACE, NOT JUST PHYSICAL SPACE***
- Optimization: could be used to find a path, or to find a controller, or parameters for a search algorithm
- In general, there are several different ways of using each algorithm
- Some of the algorithms will suck. That's OK.

What is the goal/utility?

- Mario: end of level / progress to the right edge of screen
- Pac-Man: clearing the level / score OR pills eaten OR distance to ghosts OR combination
- General Video Game Playing: winning / score OR whatever
- Fighting Game Competition: winning / health OR score

One-step search

- Start with a Game object
- Get set of moves (only four for Pac-Man)
- For each move, make a copy of Game (Game.copy())
- Try each move in its own copy (don't forget to advanceGame())
- Choose the Game state with highest score